The frontier swap: who wrote the history, who reads it

Share of next calls that repeat the history's last call. One greedy and four sampled completions per cut; no tool runs.

1. Take a recorded attempt → 2. Cut its history before call k → 3. Hand it to another model for call k only

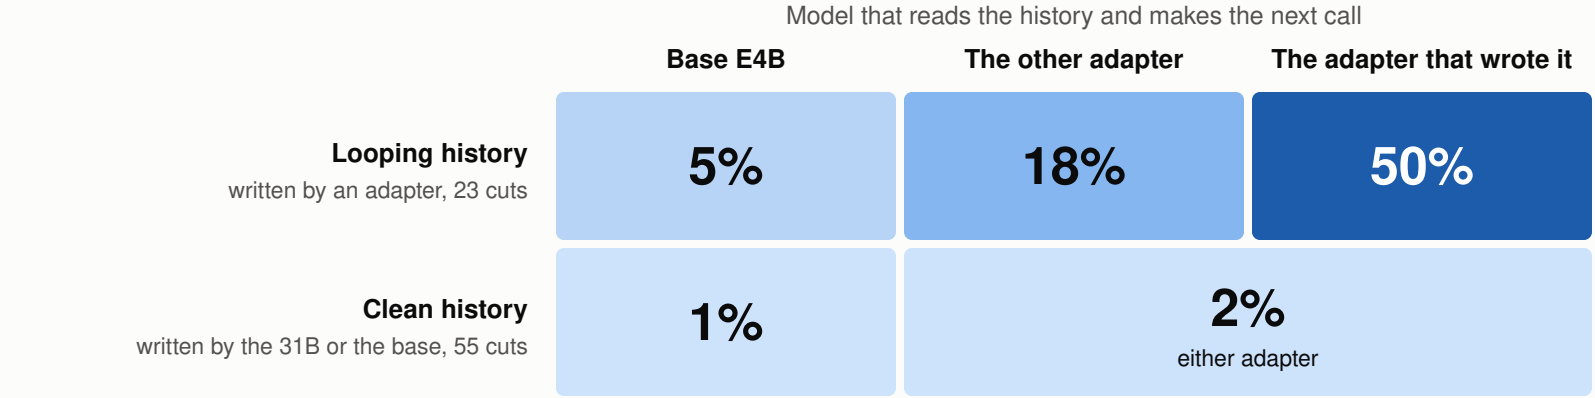

On looping histories, base minus the adapter that wrote them: −46 points (90% interval −65 to −27, task bootstrap).

The base rarely continues an adapter's loop, an adapter continues its own far more than another's, and on clean ones all repeat on 3% or less.